

```
# Imports & dataset load
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import os
from tensorflow import keras
from tensorflow.keras import layers, regularizers
from tensorflow.keras.datasets import imdb
from sklearn.metrics import confusion_matrix, classification_report

os.makedirs("images", exist_ok=True)

num_words = 10000
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num

print("Loaded:", len(train_data), "train examples and", len(test_data), "
```

Loaded: 25000 train examples and 25000 test examples

```
# Inspecting and decoding data
word_index = imdb.get_word_index()
reverse_word_index = dict((value, key) for (key, value) in word_index.items())

def decode_review(seq):
    return " ".join([reverse_word_index.get(i - 3, "?") for i in seq])

print("First review token IDs:", train_data[0][:30])
print("First label:", train_labels[0])
print("Decoded review (preview):", decode_review(train_data[0])[:300], "
```

First review token IDs: [1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65, 458, ...
First label: 1
Decoded review (preview): ? this film was just brilliant casting location

```
# Vectorizing sequences
def vectorize_sequences(sequences, dimension=num_words):
    results = np.zeros((len(sequences), dimension), dtype="float32")
    for i, sequence in enumerate(sequences):
        for j in sequence:
            if j < dimension:
                results[i, j] = 1.0
    return results

x_train = vectorize_sequences(train_data)
x_test = vectorize_sequences(test_data)
y_train = np.asarray(train_labels).astype("float32")
y_test = np.asarray(test_labels).astype("float32")
```

```
print("x_train shape:", x_train.shape)
print("y_train shape:", y_train.shape)
```

```
x_train shape: (25000, 10000)
y_train shape: (25000,)
```

```
# Building model
def build_model(hidden_layers=2, units=16, activation="tanh", output_act='
                l2_weight=0.0, dropout_rate=0.0, input_shape=(num_words,)
model = keras.Sequential()
if hidden_layers >= 1:
    model.add(layers.Dense(units, activation=activation,
                           kernel_regularizer=regularizers.l2(l2_weight),
                           input_shape=input_shape))
    if dropout_rate > 0:
        model.add(layers.Dropout(dropout_rate))
for _ in range(hidden_layers - 1):
    model.add(layers.Dense(units, activation=activation,
                           kernel_regularizer=regularizers.l2(l2_weight),
                           input_shape=input_shape))
    if dropout_rate > 0:
        model.add(layers.Dropout(dropout_rate))
model.add(layers.Dense(1, activation=output_act))
return model
```

```
# Validation split
x_val = x_train[:10000]
partial_x_train = x_train[10000:]
y_val = y_train[:10000]
partial_y_train = y_train[10000:]

print("partial_x_train shape:", partial_x_train.shape)
print("x_val shape:", x_val.shape)
```

```
partial_x_train shape: (15000, 10000)
x_val shape: (10000, 10000)
```

```
# Baseline model
baseline = build_model(hidden_layers=2, units=16, activation="tanh")
baseline.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=['accuracy'])

history = baseline.fit(partial_x_train, partial_y_train,
                      epochs=20, batch_size=512,
                      validation_data=(x_val, y_val), verbose=2)
history_baseline = history.history
```

```

/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/dense.py:93:
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
Epoch 1/20
30/30 - 3s - 103ms/step - accuracy: 0.7913 - loss: 0.4996 - val_accuracy: 0
Epoch 2/20
30/30 - 1s - 41ms/step - accuracy: 0.8981 - loss: 0.2981 - val_accuracy: 0
Epoch 3/20
30/30 - 1s - 49ms/step - accuracy: 0.9254 - loss: 0.2177 - val_accuracy: 0
Epoch 4/20
30/30 - 2s - 60ms/step - accuracy: 0.9409 - loss: 0.1703 - val_accuracy: 0
Epoch 5/20
30/30 - 2s - 63ms/step - accuracy: 0.9559 - loss: 0.1359 - val_accuracy: 0
Epoch 6/20
30/30 - 1s - 37ms/step - accuracy: 0.9646 - loss: 0.1085 - val_accuracy: 0
Epoch 7/20
30/30 - 1s - 40ms/step - accuracy: 0.9702 - loss: 0.0916 - val_accuracy: 0
Epoch 8/20
30/30 - 1s - 40ms/step - accuracy: 0.9779 - loss: 0.0710 - val_accuracy: 0
Epoch 9/20
30/30 - 1s - 37ms/step - accuracy: 0.9805 - loss: 0.0614 - val_accuracy: 0
Epoch 10/20
30/30 - 1s - 44ms/step - accuracy: 0.9808 - loss: 0.0576 - val_accuracy: 0
Epoch 11/20
30/30 - 1s - 41ms/step - accuracy: 0.9876 - loss: 0.0414 - val_accuracy: 0
Epoch 12/20
30/30 - 1s - 43ms/step - accuracy: 0.9913 - loss: 0.0340 - val_accuracy: 0
Epoch 13/20
30/30 - 2s - 57ms/step - accuracy: 0.9906 - loss: 0.0330 - val_accuracy: 0
Epoch 14/20
30/30 - 2s - 53ms/step - accuracy: 0.9988 - loss: 0.0127 - val_accuracy: 0
Epoch 15/20
30/30 - 1s - 37ms/step - accuracy: 0.9914 - loss: 0.0259 - val_accuracy: 0
Epoch 16/20
30/30 - 1s - 37ms/step - accuracy: 0.9919 - loss: 0.0267 - val_accuracy: 0
Epoch 17/20
30/30 - 1s - 38ms/step - accuracy: 0.9997 - loss: 0.0053 - val_accuracy: 0
Epoch 18/20
30/30 - 1s - 38ms/step - accuracy: 0.9935 - loss: 0.0245 - val_accuracy: 0
Epoch 19/20
30/30 - 1s - 44ms/step - accuracy: 0.9999 - loss: 0.0031 - val_accuracy: 0
Epoch 20/20
30/30 - 1s - 41ms/step - accuracy: 0.9940 - loss: 0.0201 - val_accuracy: 0

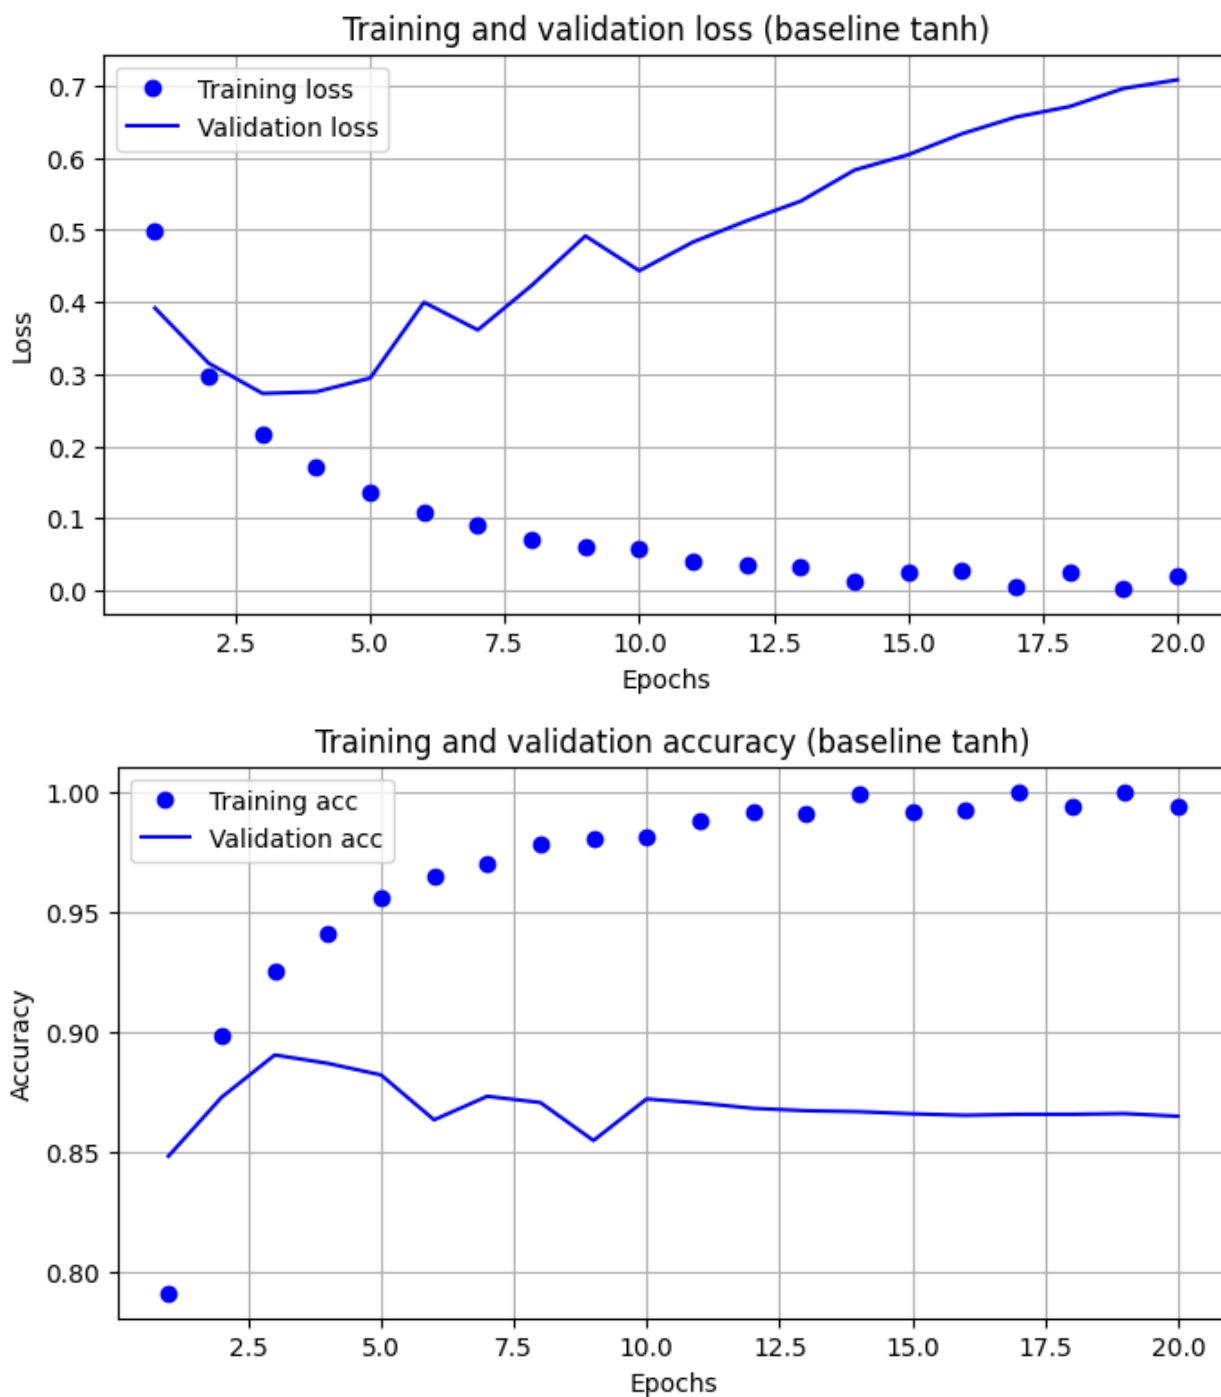
```


```

# Plotting training & validation curves
epochs = range(1, len(history_baseline["loss"]) + 1)
plt.figure(figsize=(8,4))
plt.plot(epochs, history_baseline["loss"], "bo", label="Training loss")
plt.plot(epochs, history_baseline["val_loss"], "b", label="Validation loss")
plt.title("Training and validation loss (baseline tanh)")
plt.xlabel("Epochs"); plt.ylabel("Loss"); plt.legend(); plt.grid(True)
plt.savefig("images/baseline_loss_tanh.png"); plt.show()

```

```
plt.figure(figsize=(8,4))
plt.plot(epochs, history_baseline["accuracy"], "bo", label="Training acc")
plt.plot(epochs, history_baseline["val_accuracy"], "b", label="Validation")
plt.title("Training and validation accuracy (baseline tanh)")
plt.xlabel("Epochs"); plt.ylabel("Accuracy"); plt.legend(); plt.grid(True)
plt.savefig("images/baseline_acc_tanh.png"); plt.show()
```



```
#  Retrainning final model for best epoch
best_epoch = np.argmax(history_baseline["val_accuracy"]) + 1
print("Best epoch (baseline tanh):", best_epoch)

final_model = build_model(hidden_layers=2, units=16, activation="tanh")
final_model.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])
final_model.fit(x_train, y_train, epochs=best_epoch, batch_size=512, verbose=1)

results = final_model.evaluate(x_test, y_test, verbose=2)
print("Test loss, Test accuracy:", results)
```

```
Best epoch (baseline tanh): 3
Epoch 1/3
49/49 - 3s - 63ms/step - accuracy: 0.8198 - loss: 0.4432
Epoch 2/3
49/49 - 2s - 36ms/step - accuracy: 0.9066 - loss: 0.2555
Epoch 3/3
49/49 - 2s - 36ms/step - accuracy: 0.9270 - loss: 0.1971
782/782 - 3s - 4ms/step - accuracy: 0.8859 - loss: 0.2807
Test loss, Test accuracy: [0.28068244457244873, 0.8858799934387207]
```


```
# Predictions & confusion matrix
y_prob = final_model.predict(x_test, verbose=0).flatten()
y_pred = (y_prob > 0.5).astype("int32")
cm = confusion_matrix(y_test.astype("int32"), y_pred)
print("Confusion matrix:\n", cm)
print("\nClassification report:\n", classification_report(y_test.astype("int32"), y_pred, target_names=["0", "1"]))
```

```
Confusion matrix:
[[11038  1462]
 [ 1391 11109]]
```

```
Classification report:
              precision    recall  f1-score   support

      0       0.8881      0.8830      0.8856      12500
      1       0.8837      0.8887      0.8862      12500

   accuracy          0.8859
  macro avg       0.8859      0.8859      0.8859
weighted avg       0.8859      0.8859      0.8859
```

```
#  Required experiments
configs = [
    {"id": "1layer_16", "hidden_layers": 1, "units": 16, "activation": "tanh"},
    {"id": "3layer_16", "hidden_layers": 3, "units": 16, "activation": "tanh"},
    {"id": "2layer_32", "hidden_layers": 2, "units": 32, "activation": "tanh"},
    {"id": "2layer_64", "hidden_layers": 2, "units": 64, "activation": "tanh"}]
```

```

        {"id": "2layer_mse", "hidden_layers": 2, "units": 16, "activation": "tanh"},
        {"id": "2layer_dropout", "hidden_layers": 2, "units": 16, "activation": "tanh"},
        {"id": "2layer_l2", "hidden_layers": 2, "units": 16, "activation": "tanh"}
    ]

results_list = []
for cfg in configs:
    print("\nRunning config:", cfg["id"])
    model_cfg = build_model(hidden_layers=cfg["hidden_layers"],
                            units=cfg["units"],
                            activation=cfg["activation"],
                            l2_weight=cfg["l2"],
                            dropout_rate=cfg["dropout"])
    model_cfg.compile(optimizer="rmsprop", loss=cfg["loss"], metrics=["accuracy"])
    hist = model_cfg.fit(partial_x_train, partial_y_train, epochs=10, batch_size=512,
                        validation_data=(x_val, y_val), verbose=0)
    best_val_acc = max(hist.history["val_accuracy"])
    best_epoch = np.argmax(hist.history["val_accuracy"]) + 1
    model_cfg_full = build_model(hidden_layers=cfg["hidden_layers"],
                                units=cfg["units"],
                                activation=cfg["activation"],
                                l2_weight=cfg["l2"],
                                dropout_rate=cfg["dropout"])
    model_cfg_full.compile(optimizer="rmsprop", loss=cfg["loss"], metrics=["accuracy"])
    model_cfg_full.fit(x_train, y_train, epochs=best_epoch, batch_size=512,
                      validation_data=(x_val, y_val), verbose=0)
    test_loss, test_acc = model_cfg_full.evaluate(x_test, y_test, verbose=0)
    results_list.append({
        "config": cfg["id"],
        "hidden_layers": cfg["hidden_layers"],
        "units": cfg["units"],
        "activation": cfg["activation"],
        "loss": cfg["loss"],
        "dropout": cfg["dropout"],
        "l2": cfg["l2"],
        "best_val_acc": float(best_val_acc),
        "best_epoch": int(best_epoch),
        "test_acc": float(test_acc),
        "test_loss": float(test_loss),
    })
    print("  best_val_acc=%.4f, best_epoch=%d, test_acc=%.4f" % (best_val_acc, best_epoch, test_acc))

results_df = pd.DataFrame(results_list)
results_df.to_csv("images/assignment2_results.csv", index=False)
results_df

```

```

Running config: 1layer_16
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/dense.py:93:
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
best_val_acc=0.8892 best_epoch=4 test_acc=0.8860

```